

ABHISHEK SHANKAR

774-519-0926 | Mountain View, CA | abhishekshankar88@gmail.com | linkedin.com/in/abhishekshankar97

Work Experience

RhythmX AI

December 2024 – Present

AI Engineer

Palo Alto, CA

- Built and launched production GenAI recommendation agents on Azure (App Services, AKS, Container Apps), taking LLM features from prototype to production and expanding clinical AI coverage from 4 to 200+ medical conditions at 98%+ accuracy
- Served as technical advisor to partners and stakeholders – consulting clinical, QA, and infra teams on LLM trade-offs for RAG systems, driving model selection decisions that cut API overhead by \$4,000 weekly; presented architecture and performance metrics to C-suite, accelerating decision cycles by 50%
- Engineered LLM evaluation frameworks using LLM-as-a-judge and human-in-the-loop clinical feedback – benchmarking precision, recall, and faithfulness across AI features and establishing Go/No-Go thresholds for every production deployment
- Built Batch API inference pipelines via Temporal orchestration across 200,000+ requests and profiled MongoDB queries to cut end-to-end inference latency by 50% in high-concurrency clinical workflows
- Operated in HIPAA-compliant regulated environments across HL7 v2 and FHIR data pipelines; built CI/CD pipelines for Dockerised AI services and developed reference scripts and tooling adopted company-wide across teams of varying technical depth

Boehringer Ingelheim

July 2023 – December 2023

Data Scientist Co-op

Ridgefield, CT

- Built GraphRAG knowledge extraction pipeline using Stardog and LLMs to automate NLP-driven analysis of 1,000+ supply chain PDFs, enabling structured entity extraction, text classification, and reporting at scale
- Delivered real-time analytics dashboards for 10+ drug launches using AWS CDK and Databricks, embedding ML model outputs into enterprise decision workflows for commercial teams
- Created 25+ Python and PySpark tutorial notebooks standardising ML development practices and accelerating enterprise Databricks adoption across the data science team

Adqvest Capital Pvt. Ltd.

April 2020 – June 2022

Data & Software Engineer

Chennai, India

- Built NLP sentiment classification pipelines using SpaCy and PageRank algorithms for automated text analysis, named entity recognition, and report generation, retaining 5+ clients worth \$2M+ AUM
- Deployed 100+ Python scrapers and AWS RDS infrastructure to process gigabytes of structured and unstructured data for downstream ML workflows, powering real-time market analysis for institutional clients

Projects

SupplyChainInsights – GraphRAG & LLM Fine-Tuning Platform ([GitHub](#) & [Medium](#)) May 2024 – August 2024

Python, Neo4j, LangChain, Hugging Face Transformers, OpenAI, PyTorch, RAGAS, Streamlit

- Featured in Neo4j's *This Week in Neo4j* newsletter – built a RAG platform modelling 1,000+ ARV and HIV supply chain entities with Cypher-driven traversal; fine-tuned Llama 3 and Phi-3 using Hugging Face Transformers
- Evaluated model quality using RAGAS framework measuring Faithfulness and Answer Relevancy; applied Graph Data Science similarity algorithms with OpenAI vector search to surface entity-level recommendations from structured datasets

Parallel Image Captioning with PyTorch DDP ([GitHub](#))

January – May 2024

Python, PyTorch, CUDA, Multi-GPU

- Implemented distributed data-parallel training with PyTorch DDP achieving 3x faster multi-GPU training; built encoder-decoder image captioning architecture with mixed precision training and parallel data loading pipelines

Education

Northeastern University

August 2024

Master of Science in Information Systems (Data Science)

Boston, MA

SRM Institute of Science and Technology

May 2019

Bachelor of Technology in Mechanical Engineering

Tamil Nadu, India

Skills

LLM & GenAI: Fine-tuning (Llama 3, Phi-3, GPT-4), RAG (Vector & Graph), Prompt Engineering, LLM Evaluation, LangChain, Hugging Face Transformers, OpenAI

Programming & ML: Python (Expert), PyTorch, TensorFlow, Scikit-Learn, SQL, JavaScript, PySpark

Cloud & ML Platforms: Azure ML, AWS (CDK, SageMaker, RDS), GCP, Databricks, Snowflake

Databases: Neo4j, PostgreSQL, MongoDB, Stardog (Knowledge Graphs), ChromaDB (Vector)

DevOps & Tools: RESTful APIs, CI/CD, Docker, Git, Streamlit, Airflow, Kubernetes (AKS), Tableau